



My Bookmarks

AI Bookmarks

Gitingest

Replace 'hub' with 'ingest' in any GitHub URL for a prompt-friendly text.

<https://gitingest.com/>

huggingface.co

<https://huggingface.co/spaces/mteb/leaderboard>

Vector DB Comparison

Vector DB Comparison is a free and open source tool from VectorHub to compare vector databases.

<https://superlinked.com/vector-db-comparison>

The screenshot shows a web interface titled "codebase". It has a header with the "codebase" logo and a subtitle "Turn any Git repository into a simple text digest of its codebase. This is useful for feeding a codebase into any LLM." Below the header is a form with a text input field containing "https://github.com/", a dropdown menu set to "Exclude .git, .vscode", a slider for "Include files under 50KB" (set to 50), and a checkbox for "Private Repository" which is unchecked. A large orange button labeled "Ingest" is on the right. Below the form is a link "Try these example repositories".

Papers with Code - The latest in Machine Learning

Papers With Code highlights trending Machine Learning research and the code to implement it.

<https://paperswithcode.com/>



Semantic Scholar | AI-Powered Research Tool

Semantic Scholar uses groundbreaking AI and engineering to understand the semantics of scientific literature to help Scholars discover relevant research.

<https://www.semanticscholar.org/>



Ragas

Ragas Office Hours - If you need help setting up Evals for your AI application, sign up for our Office Hours here.

<https://docs.ragas.io/en/stable/>

TruLens

Evaluation and Tracking for LLM Apps


<https://www.trulens.org/>



AI Model & API Providers Analysis | Artificial Analysis

Comparison and analysis of AI models and API hosting providers. Independent benchmarks across key performance metrics including quality, price, output speed & latency.

<https://artificialanalysis.ai/>



Quick start: Evaluating a GenAI app | Databricks Documentation

Learn how to use MLflow Evaluation to assess and improve your GenAI applications with a simple sentence completion game example.

 <https://docs.databricks.com/aws/en/mlflow3/genai/getting-started/eval>



databricks

DeepWiki | AI documentation you can talk to, for every repo

DeepWiki provides up-to-date documentation you can talk to, for every repo in the world. Think Deep Research for GitHub - powered by Devin.

 <https://deepwiki.com/>

DeepWiki

Powered by  Devin

Guardrails AI

Enforce assurance for LLM applications

 <https://www.guardrailsai.com/>

Guardrails AI

Your enterprise AI needs Guardrails.

What is Open Notebook? | Open Notebook

Take Control of Your Research, Privately

<https://www.open-notebook.ai/>

Prompt Library - Promptly AI

Access our curated collection of AI prompts. Save, organize, and reuse effective prompts for ChatGPT, Claude, Perplexity, Gemini and Deepseek. Boost your AI interactions with proven prompt templates.

 <https://www.promptly.fyi/library>

Token Economics Calculator

Tensordyne's Token Economics Calculator was developed to evaluate a broad range of AI inference systems across key metrics based on publicly available data.

 https://www.tensordyne.ai/token-economics-calculator?utm_source=newsletter&utm_medium=email&utm_campaign=token_econ_calc_nov_2025&utm_content=cta_primary&lid=1pDQEKLBRPMieODo



Chonkie

Chonkie Grounds AI In Relevant Private and Public Data

 <https://www.chonkie.ai/>



DSPy

The framework for programming—rather than prompting—language models.

 <https://dspy.ai/>

DSPy
Get Started



The framework for programming—rather than prompting—language models.


huggingface.co

<https://huggingface.co/spaces/echo840/ocrbench-leaderboard>

DeepEval - The Open-Source LLM Evaluation Framework

Native integration with Pytest, that fits right in your workflow.

🔗 <https://deepeval.com/>



Tiger Data: PostgreSQL++ for Time Series, Analytics & AI | Creators of TimescaleDB

Creators of TimescaleDB. Postgres Data Infrastructure Built for Developers, Devices, and Agents. Scale PostgreSQL effortlessly with replica sets, tiered storage, and high availability. Never pay for idle capacity. Trusted by Cloudflare, Linktree, and Polymarket.

👉 <https://www.tigerdata.com/>



SPLADE for Sparse Vector Search Explained

This article will cover the latest in learnable sparse embeddings with SPLADE — the Sparse Lexical and Expansion model.

🔗 <https://www.pinecone.io/learn/splade/>

Pinecone

LEARN

SPLADE for Sparse Vector Search Explained

AI Leaderboards 2025 - Compare All AI Models

Comprehensive AI leaderboards comparing LLM, TTS, STT, video, image, and embedding models. Compare performance, pricing, and capabilities.

🔗 <https://llm-stats.com/>

AI llm-stats.com

Analyze and compare AI models across benchmarks, pricing, and capabilities.

Spider 2.0: Enterprise Text-to-SQL Benchmark

Spider 2.0 Benchmark rigorously evaluates language models on enterprise-scale text-to-SQL tasks with complex schemas, multi-dialect SQL, and real-world workflows.

👉 <https://www.emergentmind.com/topics/spider-2-0-benchmark>

💡 Emergent Mind

LLM/Embeddings

SentenceTransformers Documentation — Sentence Transformers documentation

Sentence Transformers v4.1 just released, bringing the ONNX and OpenVINO backends to CrossEncoder (a.k.a. reranker) models. Read Cross Encoder > Usage > Speeding up Inference to read more about the performance boosts that you can expect.

🔗 <https://sbert.net/>

Nomic AI

Nomic AI is a product studio behind popular open-source AI models and the AI data infrastructure platform Atlas.

💻 <https://www.nomic.ai>

Githubs

<https://github.com/unclecode/crawl4ai>

<https://github.com/ScrapeGraphAI/Scrapetext-ai>

GitHub - mendableai/firecrawl: 🔥 Turn entire websites into LLM-ready markdown or structured data. Scrape, crawl and extract with a single API. - mendableai/firecrawl
🔥 Turn entire websites into LLM-ready markdown or structured data. Scrape, crawl and extract with a single API. - mendableai/firecrawl
🔗 <https://github.com/mendableai/firecrawl?tab=readme-ov-file>

<https://github.com/docling-project/docling>

<https://github.com/QuivrHQ/MegaParse>

<https://github.com/enoch3712/ExtractThinker>

<https://github.com/Giskard-AI/giskard>

<https://github.com/explodinggradients/ragas>

<https://github.com/truera/trulens/>

<https://github.com/landing-ai/agentic-doc>

<https://github.com/guardrails-ai/guardrails>

<https://github.com/Shubhamsaboo/awesome-llm-apps>

https://github.com/NirDiamant/RAG_Techniques

<https://github.com/athina-ai/rag-cookbooks>

<https://github.com/FareedKhan-dev/all-rag-techniques>

<https://github.com/guyernest/advanced-rag>

<https://github.com/NirDiamant/agents-towards-production>

<https://github.com/ashishpatel26/500-AI-Agents-Projects>

GitHub - aurelio-labs/semantic-router: Superfast AI decision making and intelligent processing of multi-modal data.

Superfast AI decision making and intelligent processing of multi-modal data. - aurelio-labs/semantic-router

<https://github.com/aurelio-labs/semantic-router/tree/main>

aurelio-labs/**semantic-router**

Superfast AI decision making and intelligent processing of multi-modal data.

45 Contributors 496 Used by 25 Discuss

<https://github.com/DrewThomasson/ebook2audiobook>

https://github.com/jingyaogong/minimind/blob/master/README_en.md

<https://github.com/Tencent-Hunyuan/HunyuanOCR>

<https://github.com/deepseek-ai/DeepSeek-OCR>

<https://github.com/PaddlePaddle/PaddleOCR>

Python Bookmarks

click

Composable command line interface toolkit

<https://pypi.org/project/click/>



rich

Render rich text, tables, progress bars, syntax highlighting, markdown and more to the terminal

 <https://pypi.org/project/rich/>



Client Challenge

 <https://pypi.org/project/extract-thinker/0.0.1/>

colab-xterm

Open a terminal in colab, including the free tier.

 <https://pypi.org/project/colab-xterm/>



Database Bookmarks

Milvus | High-Performance Vector Database Built for Scale

Milvus is an open-source vector database built for GenAI applications. Install with pip, perform high-speed searches, and scale to tens of billions of vectors.

 <https://milvus.io/>


Vector database built
for scalable similarity search

 milvus

Chroma

Chroma is the open-source AI application database. Batteries included.

 <https://www.trychroma.com/>



Qdrant - Vector Database

Qdrant is an Open-Source Vector Database and Vector Search Engine written in Rust. It provides fast and scalable vector similarity search service with convenient API.

 <https://qdrant.tech/>

 Qdrant
Qdrant Vector Database

The AI-native database developers love | Weaviate

Bring AI-native applications to life with less hallucination, data leakage, and vendor lock-in

 <https://weaviate.io/>



The vector database to build knowledgeable AI | Pinecone

Search through billions of items for similar matches to any object, in milliseconds. It's the next generation of search, an API call away.

 <https://www.pinecone.io/>


 Pinecone
The vector database for
scale in production.

Cloud Bookmarks

How to monitor, optimize, and secure Amazon Cognito machine-to-machine authorization | Amazon Web Services

Amazon Cognito is a developer-centric and security-focused customer identity and access management (CIAM) service that simplifies the process of adding user sign-up, sign-in, and access control to your mobile and web applications. Cognito is a highly available service that supports a range of use cases, from managing user authentication and authorization to enabling

 <https://aws.amazon.com/blogs/security/how-to-monitor-optimize-and-secure-amazon-cognito-machine-to-machine-authentication/>



Rent GPUs | Vast.ai

Rent high-performance cloud GPUs at low cost with Vast.ai. Instantly deploy GPU rentals for AI, machine learning, deep learning, and rendering. Flexible pricing, fast setup, and global availability

 <https://vast.ai/>

vast.ai

Utility Bookmarks

Excalidraw to PDF converter

 <https://luis-c465.github.io/excalidraw-pdf-splitter/>

Material Design Icons - Icon Library - Pictogrammers

The original. Following Google's Material Design guidelines for system icons, MDI is our largest library, touting over 7200 unique icons!

 <https://pictogrammers.com/library mdi/>



Pictogrammers

Open-source iconography for
designers and developers

SvgPathEditor

Online editor to create and manipulate SVG paths

 <https://yqnn.github.io/svg-path-editor/>

Speech and Language Processing

 <https://web.stanford.edu/~jurafsky/slp3/>